



DANI CABALLERO MÉLICH

AI ENGINEER · AGENTIC SYSTEMS

Barcelona · remote-first
EU citizen, no sponsorship needed
dani.caballero.m@gmail.com
+34 644 974 715
linkedin.com/in/dani-caballero-
melich
github.com/DaniMelich ·
danimelich.com

STACK

Agents LLM agents · LangGraph ·
MCP server · Pydantic contracts ·
multi-agent orchestration · tool
use · prompt engineering ·
human-in-the-loop gates

Evaluation labeled eval sets ·
accuracy gates · ground-truth
verification · CI gates ·
observability

RAG / retrieval hybrid dense +
lexical fusion · embeddings ·
gold-set scoring · recall@k · MRR

Cost token budgets · per-stage
model and effort routing ·
fallback chains

Providers Anthropic · OpenAI ·
Claude Code

PROGRAMMING

Python 3.12 · C# · Kotlin · Rust ·
TypeScript · SQLite · asyncio ·
pytest

ALSO

Self-taught from seventeen: 15+
prototypes in six years before
turning professional. Fully remote
throughout, cross-timezone with
a US studio. Catalan and Spanish
native · English fluent.

I build multi-agent LLM systems that run unattended and merge working code into a commercial Unity game overnight. I price what agents cost, eval what they claim, and cut what the numbers kill. Two years on AI systems, on top of shipped product work.

AGENT SYSTEMS

Dungeon Heroes pipeline and 16-agent fleet

2026 · self-directed · the platform that builds my mobile title — signed release builds, store-ready

- **Sixteen agents ship my game while I sleep.** I rebuilt my agent platform Claude-native as a fleet of 16 agents writing production code into the game overnight: model tier picked by cost of failure, worktree lanes derived from co-change analysis of the repo's history, and one fixed dispatch contract for every agent — branch, files touched, what to verify, confidence. It merged 107 commits in its first full day, unattended.
- **Before the fleet, I built the brief-to-code pipeline** — 20 LangGraph stages, typed Pydantic contracts, 5 human gates, exposed as an MCP server — and I made it price its own work: \$67–98 per human-accepted screen, every rerun charged to its cause.
- **Then I killed it on its own telemetry.** The ledger showed the validator layer buying nothing — zero true positives across 730 findings in three audited runs, close to half of all spend going to reruns — so I archived all 35,000 lines and moved verification onto ground truth: compiler, filesystem, live scene. The fleet above is its replacement, and it ships more.
- **I cut the expensive share of my model bill from 42.0% to 12.1%,** against a 20% bar: a token-economics program over 30 days of my own usage — spend measured per context bucket, levers ranked by ROI, the top two shipped and re-measured.
- **I wrote the eval before the model got the job.** I gated a local-model swap on an 80-item labeled intent eval with a 0.95 accuracy bar; both candidates failed, at 77.5% and 85.0%, and neither shipped.
- **I scored retrieval before any stage relied on it:** hybrid dense and lexical retrieval fused by reciprocal rank, scored on a 32-pair gold set built from real failures — recall@10 0.969.

EXPERIENCE

Somni Game Studios · Software Engineer

2025 – 2026 · Stepland, live on Google Play and the App Store · one of three core programmers in a larger team

- **I landed the rebuilt tutorial on the opening day of the studio's 4YFN showcase** — 82 declarative steps behind a sequence runner with resume and telemetry — by driving coding-agent sessions inside the production repository.

Purple Pwny Studios · Software Engineer

2023 – 2024 · remote, US studio

- **I gave a real-money wagering game a skilled player to balance against:** a hand-authored heuristic policy agent run through Monte Carlo simulation, so the economy was tuned against strong play instead of random rolls.